A Scanner-Informed Synthetic Overlay Benchmark Dataset for EUV Lithography Calibration

Aris Magklaras*, George Tsirogiannis †. Alexios Birbas*, Christos Gogos‡, Panayiotis Alefragis†
*Electrical and Computer Engineering Department, University of Patras, Rio, Greece
†University of Patras, Patras, Greece
‡Informatics and Telecommunications Department, University of Ioannina, Ioannina, Greece
*a.magklaras@upnet.gr, †alefrag@uop.gr, ‡cgogos@uoi.gr,
*birbas@ece.upatras.gr, *gtsirogianni@upatras.gr

***Abstract*—We present EUV-OVL-SYN, an open and fully reproducible synthetic benchmark dataset for developing and evaluating overlay correction algorithms in 300 mm wafer extreme-ultraviolet (EUV) lithography. The dataset preserves the physical signal hierarchy of real scanner data, namely, a dominant inter-field Brink/van den Brink polynomial fingerprint (cross-wafer $R^2 \geq 0.93$), four distinct lot-level temporal drift scenarios modelling linear scanner heating, reticle alignment drift, multi-parameter process drift, and chiller-cycling oscillation, and a nonlinear intra-field residual analytically verified to lie outside the Brink polynomial basis. The dataset comprises four lots of 50 wafers each at 1 420 metrology targets each ($284{,}000$ measurements total). A companion two-stage correction experiment demonstrates that a Physics-Guided Neural Network (PGNN) residual model achieves a 33–52% reduction in overlay 3σ over the Brink-only baseline, validating the dataset's utility as a rigorous benchmark.**

***Index Terms*—EUV lithography, overlay error, synthetic benchmark, Brink model, physics-informed neural networks, advanced process control, semiconductor metrology, intra-field variation.**

I. Introduction

Overlay is the alignment accuracy between consecutive patterning layers and it is among the most critical manufacturing and monitoring parameters governing transistor density and yield at advanced nodes. EUV lithography, especially at sub-5 nm nodes requires extremely tight overlay budgets [1], [2], with the transition to High-NA EUV (0.55-NA) pushing overlay requirements below 1 nm [3]. Achieving this demands accurate correction of the systematic overlay fingerprint induced by scanner (i.e. EUV photolithographic machines) optics, stage mechanics, and chuck thermal cycling.

The classical correction tool is the Brink/van den Brink polynomial model [4], [5], which expresses the inter-field overlay error as a 22-parameter weighted combination of wafer-stage coordinates. Fitted once per lot based on a limited set of training wafers, it reduces overlay 3σ metric from $\sim$7–9 nm (raw) to $\sim$1.5–2.5 nm [6]. The remaining systematic residual, arising from thermal related asymmetries [7], reticle bowing, and intra-field aberrations [8], [9] — cannot be described by any fixed polynomial of this clas of models, and lately, machine learning (ML) post-correction techniques have been proposed [10]–[13].

A limiting barrier to ML research in this domain is the absence of open and realistic datasets [14]. Real scanner overlay data are commercially sensitive; their publication might violate confidentiality agreements between chipmakers and scanner OEMs. Existing public simulation datasets either model overlay as isotropic Gaussian noise (providing no systematic structure for learning) [8] or use simplified 6–10 parameter field models [10]–[12] that do not replicate the separability of the Brink fingerprint from a nonlinear learnable residual.

This paper addresses this gap. It proposes a synthetic dataset that is designed from "first principles" (in terms of physical contributors) to be maximally useful as a benchmarking substrate. The key contributions are:

1) A physically grounded signal model replicating the scale, structure, and wafer-to-wafer correlation of real production EUV overlay data.
2) A parameter-calibrated temporal drift library covering four qualitatively distinct lot-level drift mechanisms.
3) A nonlinear intra-field residual whose functional form is analytically disjoint from the Brink basis, thus providing a clean, controllable learning target (orthogonal by design to the remaining systematic residuals).
4) Full reproducibility: a single Python script regenerates the complete dataset from a fixed seed, that is easy to understand and extend.

II. Related Work

A. Overlay Error Modeling

The Brink polynomial [4] expresses overlay as a linear combination of inter-field stage coordinates and intra-field die offsets, extended by van den Brink *et al.* [5] to include 3rd- and 5th-order radial distortion terms. Fitting 22 parameters by weighted least squares on a set of training wafers captures $> 90\%$ of the systematic overlay variance at current EUV technology nodes. Alternative physical models based on Zernike polynomial decomposition of the scanner pupil [9], [15] provide a complementary description of the intra-field aberration contribution, as demonstrated by [8] for photolithography simulation. Our work is complementary and extends [8], that focuses on the aberration-to-overlay transfer within a single exposure field, EUV-OVL-SYN models the complete lot-level hierarchy including cross-wafer temporal drift.

B. Machine Learning for Overlay Correction

Data-driven correction of Brink residuals has been explored using Gaussian process regression, random forests, and deep neural networks [10]–[12], [16]. Physics-guided neural networks (PGNNs) [13], [17] have attracted attention for their ability to preserve spatial patterns and incorporate rotational symmetry priors that are known to govern overlay distortion fields. All of these methods require an annotated dataset with known ground-truth decomposition to validate that a model corrects the Brink residual rather than overfitting a specific Brink fingerprint, a requirement that EUV-OVL-SYN is designed to satisfy.

III. Dataset Architecture

A. Signal Decomposition Model

The observed overlay vector at site i on wafer w is modelled as

$$\mathrm{OVL}_{w,i} = \mathbf{H}_i\,\boldsymbol{\beta}_w + \mathbf{f}_{\mathrm{NL}}(x_i, y_i, w) + \boldsymbol{\varepsilon}_{w,i}, \tag{1}$$

where $\mathrm{OVL}_{w,i} \in \mathbb{R}^2$ is the (X, Y) overlay pair at site i, $\mathbf{H}_i \in \mathbb{R}^{2\times 22}$ is the block-diagonal Brink design matrix at site i (§ III-C), $\boldsymbol{\beta}_w = \boldsymbol{\beta}_{\mathrm{lot}} + \Delta\boldsymbol{\beta}(w) \in \mathbb{R}^{22}$ is the per-wafer Brink parameter vector stacking 11 X-axis and 11 Y-axis coefficients (§ III-D), $\mathbf{f}_{\mathrm{NL}} \in \mathbb{R}^2$ is the nonlinear intra-field residual (§ III-E), and $\boldsymbol{\varepsilon}_{w,i} \sim \mathcal{N}(\mathbf{0},\, \sigma^2\mathbf{I}_2)$ with $\sigma = 0.15$ nm is isotropic Gaussian metrology noise applied independently to the X and Y channels. All three components are generated independently and with full ground-truth access, enabling exact decomposition of any prediction error.

B. Scanner and Wafer Geometry

EUV-OVL-SYN follows the spatial structure of an ASML NXE-class 300 mm scanner wafers. Exposure fields are arranged on a rectangular grid with pitch 26×33 mm; only fields whose centre lies within a 138 mm keep-out radius are retained, resulting **71 fields per wafer** (Fig. 1). The wafer coordinates are normalised by the nominal radius (i.e.)$R = 300$ mm) for the Brink model ($x^* = X_w/R \le 0.46$), consistent with the original formulation [5].

Within each exposure field, a typical 4×5 intra-field die grid is placed at die offsets $x_d \in \{-9, -3, 3, 9\}$ mm, $y_d \in \{-12, -6, 0, 6, 12\}$ mm, giving **20 sites per field** and $71\times 20 =$ **1 420** sites per wafer. The full dataset contains $4 \times 50 \times 1{,}420 =$ **284,000** overlay targets across the six-column schema $(X_w, Y_w, x_d, y_d, \mathrm{OVL}_X, \mathrm{OVL}_Y)$, directly compatible with real APC (Automatic Process Control) metrology exports.

C. Extended Brink/van den Brink Model

Stage-level overlay is modelled by the 22-parameter Brink polynomial [4], [5]. For each site the design matrix rows are

$$\mathbf{h}_X = [1,\ x^*,\ -y^*,\ (y^*)^2,\ x,\ y,\ -x^2,\ -xy,\ y^2,\ xr^2,\ xr^4],$$
$$\mathbf{h}_Y = [1,\ -x^*,\ y^*,\ (x^*)^2,\ y,\ x,\ -y^2,\ -yx,\ x^2,\ yr^2,\ yr^4],$$

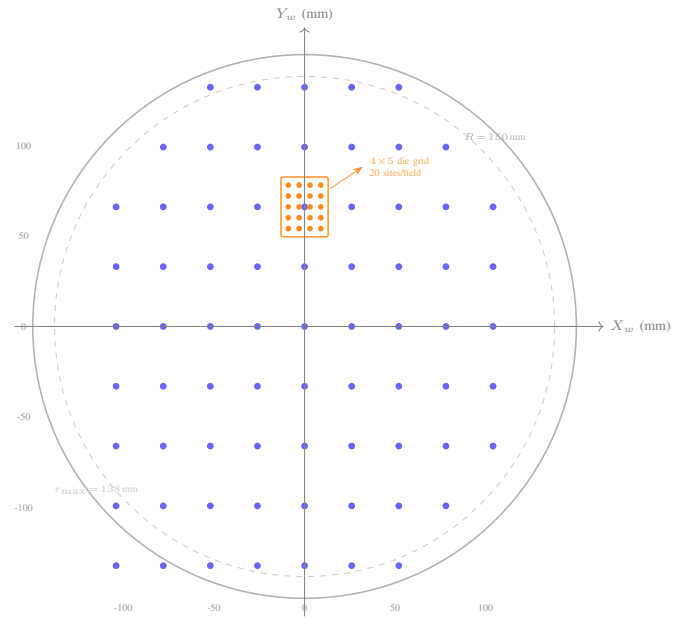

Fig. 1. Exposure-field grid for EUV-OVL-SYN (71 fields, •) on a 300 mm wafer. The dashed circle marks the $r = 138$ mm keep-out boundary. One field (orange box) is expanded to show the 4×5 intra-field die grid (20 sites).

TABLE I
Dominant Brink Baseline Parameters (Representative Values)

Param.	Meaning	Value	Gain†
T_x	Inter-field translation	1.50 nm	1
M_x	Inter-field magnification	7.50 nm	$\lvert x^*\rvert \le 0.46$
R_x	Inter-field rotation	1.20 nm	$\lvert y^*\rvert \le 0.46$
m_x	Intra-field magnification	0.070 nm/mm	$\lvert x\rvert \le 9$ mm
D_{3x}	3rd-order radial	8×10^{-5} nm/mm^3	xr^2
D_{5x}	5th-order radial	2×10^{-9} nm/mm^5	xr^4

† Max. column value in $\mathbf{H}$ across all sites.

where $(x^*, y^*) = (X_w, Y_w)/R$ are the normalised inter-field coordinates and $(x, y) \equiv (x_d, y_d)$ are the intra-field die offsets. The 11 parameters per axis capture global translation, magnification, rotation, intra-field bilinear terms, and 3rd/5th-order radial distortion. A common lot-level baseline $\boldsymbol{\beta}_{\mathrm{lot}}$ (Table I) receives a deterministic, lot-specific fixed offset $\delta\boldsymbol{\beta}_{\mathrm{lot}}$ applied to a small subset of dominant parameters — T_x, M_x, T_y, M_y, and at most one additional parameter per lot from $\{R_x,\ B_x,\ B_y\}$ — with amplitudes in the range $[-1.1,\ +1.2]$ nm, to represent different scanner set-points across lots. Lot 1 serves as the reference (zero offset); Lots 2–4 carry non-zero offsets on four to five parameters.

D. Lot-Level Temporal Drift Scenarios

Real scanner lots exhibit systematic wafer-to-wafer parameter evolution driven by lens heating [7], reticle thermal expansion, and chiller cycling. EUV-OVL-SYN provides four calibrated scenarios:

- **Lot 1 — Linear scanner heating:** T_x drifts $+1.5$ nm, M_x drifts -0.6 nm over the lot, consistent with documented lens warm-up profiles in immersion and EUV scanners [1], [7].
- **Lot 2 — Reticle alignment drift:** T_y creep $+1.8$ nm, M_y drift -0.8 nm, mimicking slow thermal expansion of the reticle stage clamp.
- **Lot 3 — Multi-parameter drift:** Six Brink parameters drift simultaneously (largest amplitude), exercising the correction model under the most challenging lot-level nonstationarity.
- **Lot 4 — Chiller-cycling oscillation:** T_x and T_y oscillate sinusoidally with ~ 2.5 cycles over the lot (amplitude 0.9 nm), representing the periodic thermal perturbation of an imperfectly regulated chiller.

All four scenarios include per-parameter iid process jitter $\delta\boldsymbol{\beta}_w \sim \mathcal{N}(\mathbf{0}, \mathrm{diag}(\boldsymbol{\sigma}_\beta^2))$, with per-parameter scales $\boldsymbol{\sigma}_\beta$ calibrated so that each parameter contributes ≈ 0.05 nm of wafer-to-wafer overlay noise (i.e., proportional to the inverse effective column gain of $\mathbf{H}$). This calibration prevents high-gain 5th-order terms (D_{5x}: gain $\sim 10^5$ mm^5) from dominating the noise budget, a physically motivated constraint absent from naive uniform-perturbation approaches.

E. Nonlinear Intra-Field Residual Design

The PGNN learning target f_{NL} must satisfy two properties:

1) It depends on intra-field die coordinates (x_d, y_d) only — matching the PGNN input domain — not on inter-field stage coordinates.
2) Its functional form lies outside the Brink polynomial column space, so the Brink model cannot capture it.

We verified analytically that the following terms satisfy both requirements. For polynomial terms, verification amounts to checking that the target monomial is absent from the Brink column space of the relevant axis; for the sinusoidal term, transcendental functions cannot be represented by any finite polynomial, so membership is excluded by construction: (a) $\sin(\pi\tilde{x})\cos(\frac{1}{2}\pi\tilde{y})$ — sinusoidal in die normalised coordinates $(\tilde{x}, \tilde{y}) = (x_d/9,\, y_d/12)$, outside any polynomial basis; (b) $\tilde{x}^2\tilde{y}$ — not in the Brink X-equation, whose highest-degree intra-field column is $xr^2 = x^3 + xy^2$, which does not contain the x^2y monomial; (c) $\tilde{x}\tilde{y}^2$ — not in the Brink Y-equation, whose corresponding column is $yr^2 = x^2y + y^3$, which does not contain xy^2.

The complete residual is

$$f_{\mathrm{NL},X} = A_\ell\,(1 + 0.3\tilde{w})\big[1.8\sin(\pi\tilde{x})\cos(\tfrac{\pi\tilde{y}}{2}) + 1.2\,\tilde{x}^2\tilde{y}\big], \quad (2)$$

$$f_{\mathrm{NL},Y} = A_\ell\,(1 + 0.3\tilde{w})\big[1.5\sin(\pi\tilde{y})\cos(\tfrac{\pi\tilde{x}}{2}) + 1.0\,\tilde{x}\tilde{y}^2\big], \quad (3)$$

where $A_\ell \in \{1.0, 0.80, 1.30, 0.95\}$ is the lot-specific amplitude, $\tilde{w} = (w-1)/(W-1)$ is the normalised wafer index, and the factor $(1 + 0.3\tilde{w})$ encodes a 30% amplitude growth mimicking slow lens-heating accumulation. The resulting nonlinear residual has $1\sigma \approx 1.1$–1.4 nm per axis, well above the 0.15 nm noise floor — providing a clear, controllable learning target (Fig. 4).

TABLE II
PER-LOT OVERLAY STATISTICS ($N_{\mathrm{train}} = 10$, ALL TEST WAFERS)

Lot	**Drift Type**	**Raw** 3σ (nm)	**Brink** 3σ (nm)	R^2	**P95** (nm)
1	Linear (scanner)	8.04	1.782	0.956	1.095
2	Linear (reticle)	7.48	1.532	0.961	0.973
3	Multi-param.	8.89	2.207	0.938	1.350
4	Sinusoidal	8.05	2.578	0.897	1.822
Mean		8.12	2.025	0.938	1.310

F. Metrology Noise

Site-level gaussian noise $\varepsilon_{w,i} \sim \mathcal{N}(0, 0.0225)$ ($\sigma = 0.15$ nm) is added independently to every X and Y measurement, consistent with reported single-measurement repeatability of current EUV metrology tools [6]. The signal-to-noise ratio (SNR) of the nonlinear residual relative to measurement noise is approximately 7–9 dB, ensuring that the learning signal is detectable but not trivially dominant.

IV. DATASET STATISTICS AND CHARACTERISATION

Table II summarises the per-lot overlay statistics measured on held-out test wafers (wafers 13–50 after a 10-wafer training, 2-wafer validation split). The raw overlay 3σ ranges from 7.5 to 8.9 nm across lots, consistent with values reported for pre-corrected EUV overlay in production-like conditions [6]. After applying the cross-wafer Brink model, 3σ drops to 1.5–2.6 nm, confirming that the synthetic Brink fingerprint ($R^2 \geq 0.93$) has the expressive power expected of a realistic scanner dataset. The remaining structured residual (Brink 3σ) directly corresponds to f_{NL}, since the noise contribution is only $3\sigma_\varepsilon \approx 0.45$ nm.

Figure 2 shows the overlay decomposition for Lot 1, Wafer 25 (mid-lot). The raw field (a) exhibits a smooth, rotationally-structured pattern dominated by the Brink fingerprint. The Brink prediction (b) captures this pattern closely, and the Brink residual (c) reveals the underlying nonlinear structure. Panel (d) shows the ground-truth nonlinear component computed directly from (2)–(3), confirming that the residual is a deterministic, spatially structured target rather than noise.

Figure 3 illustrates the wafer-to-wafer evolution of the per-wafer mean overlay and 3σ across all 50 wafers for each lot. The four qualitatively distinct drift fingerprints are clearly visible: monotonic creep in Lots 1–2, accelerating multi-parameter drift in Lot 3, and the characteristic oscillatory signature of chiller cycling in Lot 4. Polynomial thermal detrending (degree $K = 2$) removes these trends before Brink fitting, reducing the effective noise seen by the model.

The spatial structure of the nonlinear residual across all four lots is shown in Fig. 4. The pattern is fixed to the intra-field die grid, grows gently over the lot (30% amplitude increase), and shows distinct orientation and amplitude per lot. Lot 3 has the largest amplitude ($A_3 = 1.30$), consistent with its strongest Brink residual in Table II.

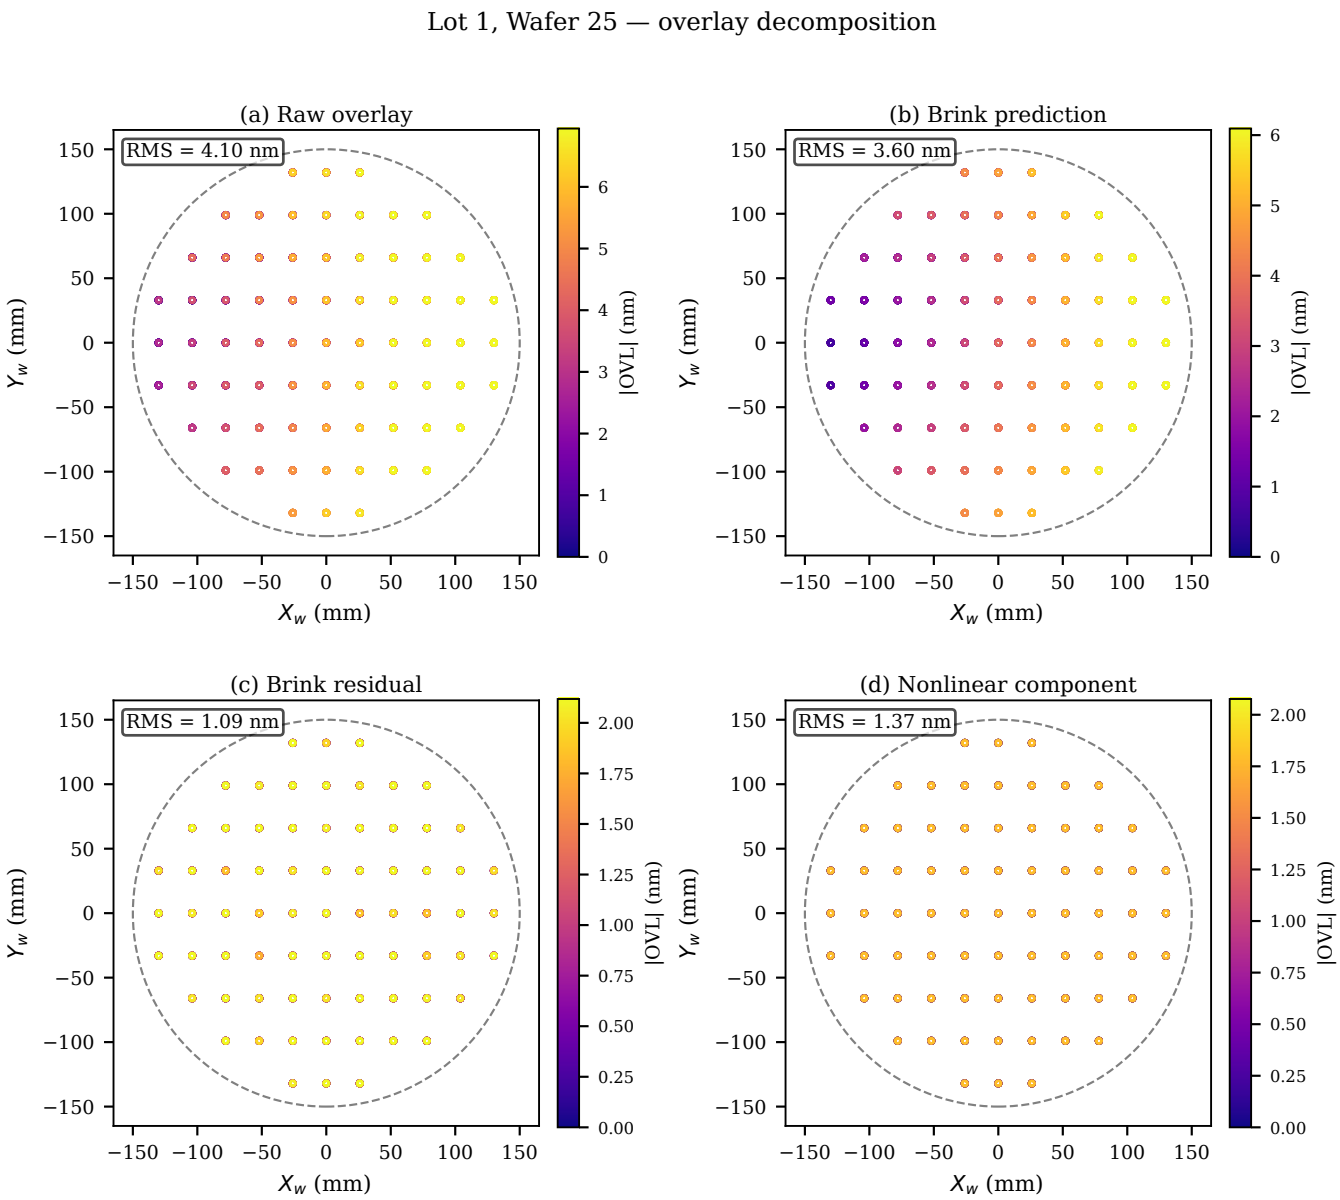

Fig. 2. Overlay decomposition for Lot 1, Wafer 25. **(a)** Raw overlay field ($3\sigma \approx 3.8$ nm). **(b)** Brink model prediction (cross-wafer fit, $N_{\text{train}} = 10$). **(c)** Brink residual revealing the nonlinear intra-field structure ($3\sigma \approx 1.8$ nm). **(d)** Ground-truth nonlinear component f_{NL} — the intended ML learning target. Colour encodes overlay magnitude (nm); arrows indicate direction.

V. Benchmark Experiments

To validate the proposed dataset, we evaluate a two-stage correction pipeline, similar to the one introduced in [8] but operating on overlay rather than optical (i.e. Zernike) coefficients. comprising: (i) a cross-wafer Brink model (Stage 1), and (ii) a Physics-Guided Neural Network (PGNN) residual correction fitted on the Brink residuals of $N_{\text{train}} \in \{3, 5, 10\}$ training wafers [13]. The PGNN is a 3-hidden-layer MLP (Multi Layer Perceptron [18]) with $\tanh$ activations taking normalised die coordinates $(\tilde{x}, \tilde{y})$ as inputs, regularised by a spatial smoothness penalty on the gradient $\|\nabla_{(\tilde{x},\tilde{y})} f\|^2$.

Table III reports the overlay 3σ on held-out test wafers across all four lots and training set sizes. The PGNN achieves a **33–52%** reduction in Brink 3σ at $N_{\text{train}} = 10$, with a 100% per-wafer improvement rate across all conditions. Importantly, the idela improvement is achieved with minimum traing, i.e. **near-constant** with respect to training set size: as few as **three training wafers** suffice to identify the nonlinear fingerprint, a property that could lead to a future direct real-time APC deployment at lot start.

The lower improvement on Lot 4 (15.4%) is expected: the sinusoidal chiller-cycling drift induces a wafer-to-wafer variation in the Brink parameters (T_x, T_y) that, after polynomial detrending, leaves a residual that is partially time-varying. This

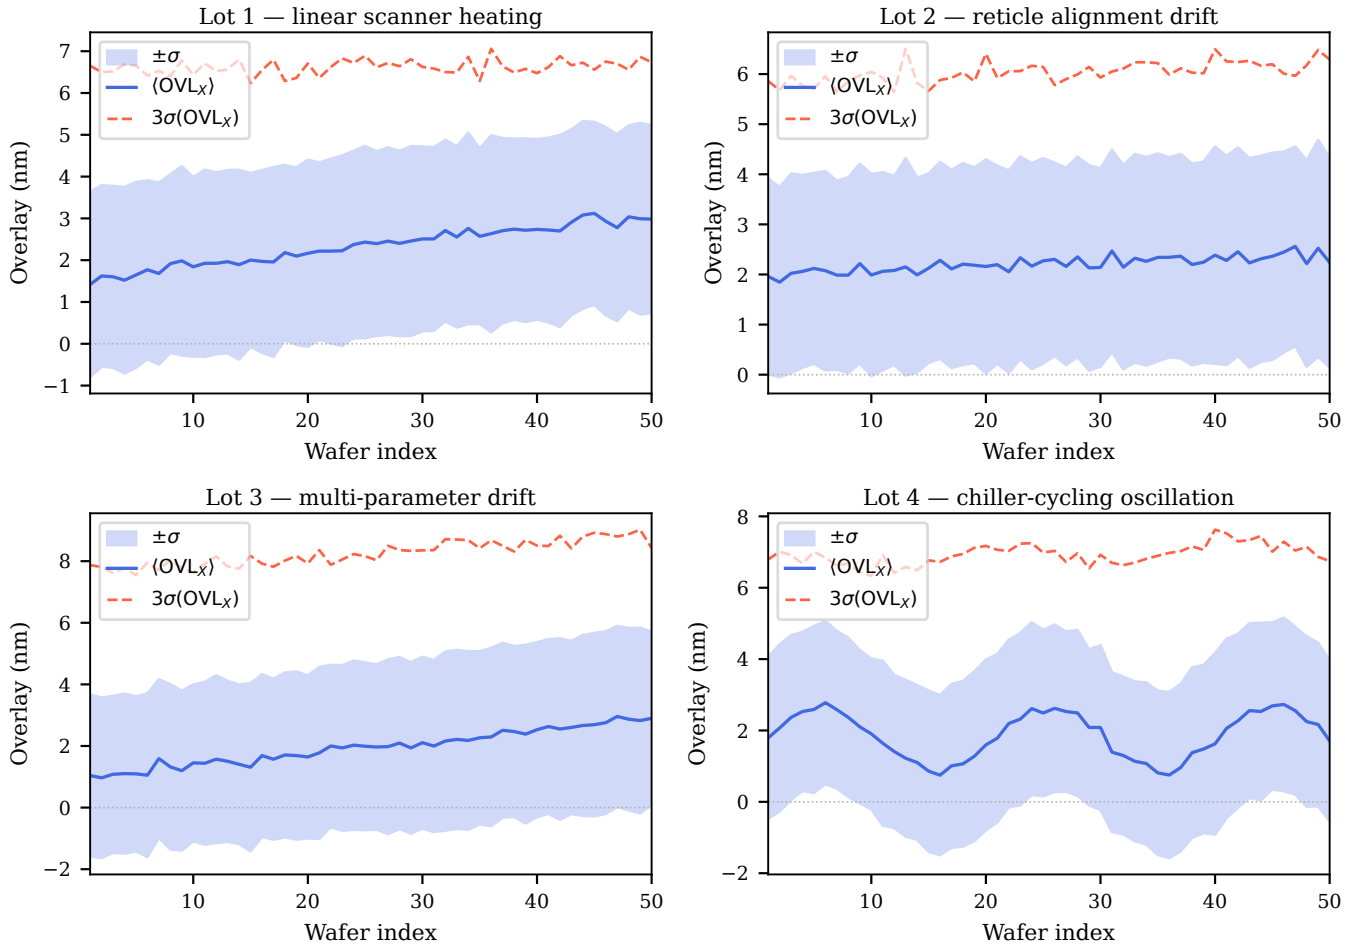

Fig. 3. Per-wafer temporal drift profiles across all 50 wafers in each lot. Blue solid: per-wafer mean OVL_X. Red dashed: per-wafer $3\sigma(\mathrm{OVL}_X)$. The four drift scenarios (linear heating, reticle drift, multi-parameter drift, chiller oscillation) are clearly distinguishable and span the primary mechanisms reported in EUV production environments [6], [7].

TABLE III
BENCHMARK: BRINK 3σ VS. BRINK+PGNN 3σ (NM)

Lot	**Brink**	**Brink+PGNN** (N_{train})			**Improv.**
		$n=3$	$n=5$	$n=10$	($n=10$)
1	1.782	1.024	1.012	1.005	43.6%
2	1.532	1.031	0.998	0.976	36.3%
3	2.207	1.142	1.112	1.068	51.6%
4	2.578	2.527	2.126	2.181	15.4%
Mean					**36.7%**

makes Lot 4 a harder benchmark scenario, and its distinct character from Lots 1–3 is an intentional feature of the dataset design.

VI. Conclusion

We propose a realistic, fully reproducible synthetic benchmark dataset for overlay correction research in EUV mass production lithography. The dataset faithfully replicates the signal hierarchy of real scanner data, a dominant Brink polynomial fingerprint, calibrated temporal drift in four production-relevant scenarios, and a nonlinear intra-field residual analytically disjoint from the Brink polynomial basis, while providing full ground-truth decomposition unavailable from industrial sources. A companion benchmark demonstrates that a PGNN two-stage correction pipeline achieves 33–52% reduction in overlay 3σ over the Brink-only baseline with a 100% per-wafer win rate, confirming that the dataset is both non-trivially challenging and solvable. EUV-OVL-SYN is intended as an open access industrial standard dataset for the semiconductor metrology community to develop, compare, and rigorously validate the next generation of overlay correction algorithms.

Acknowledgements

This work was partially supported through the HCCC Chips-JU project (GA 101217803).

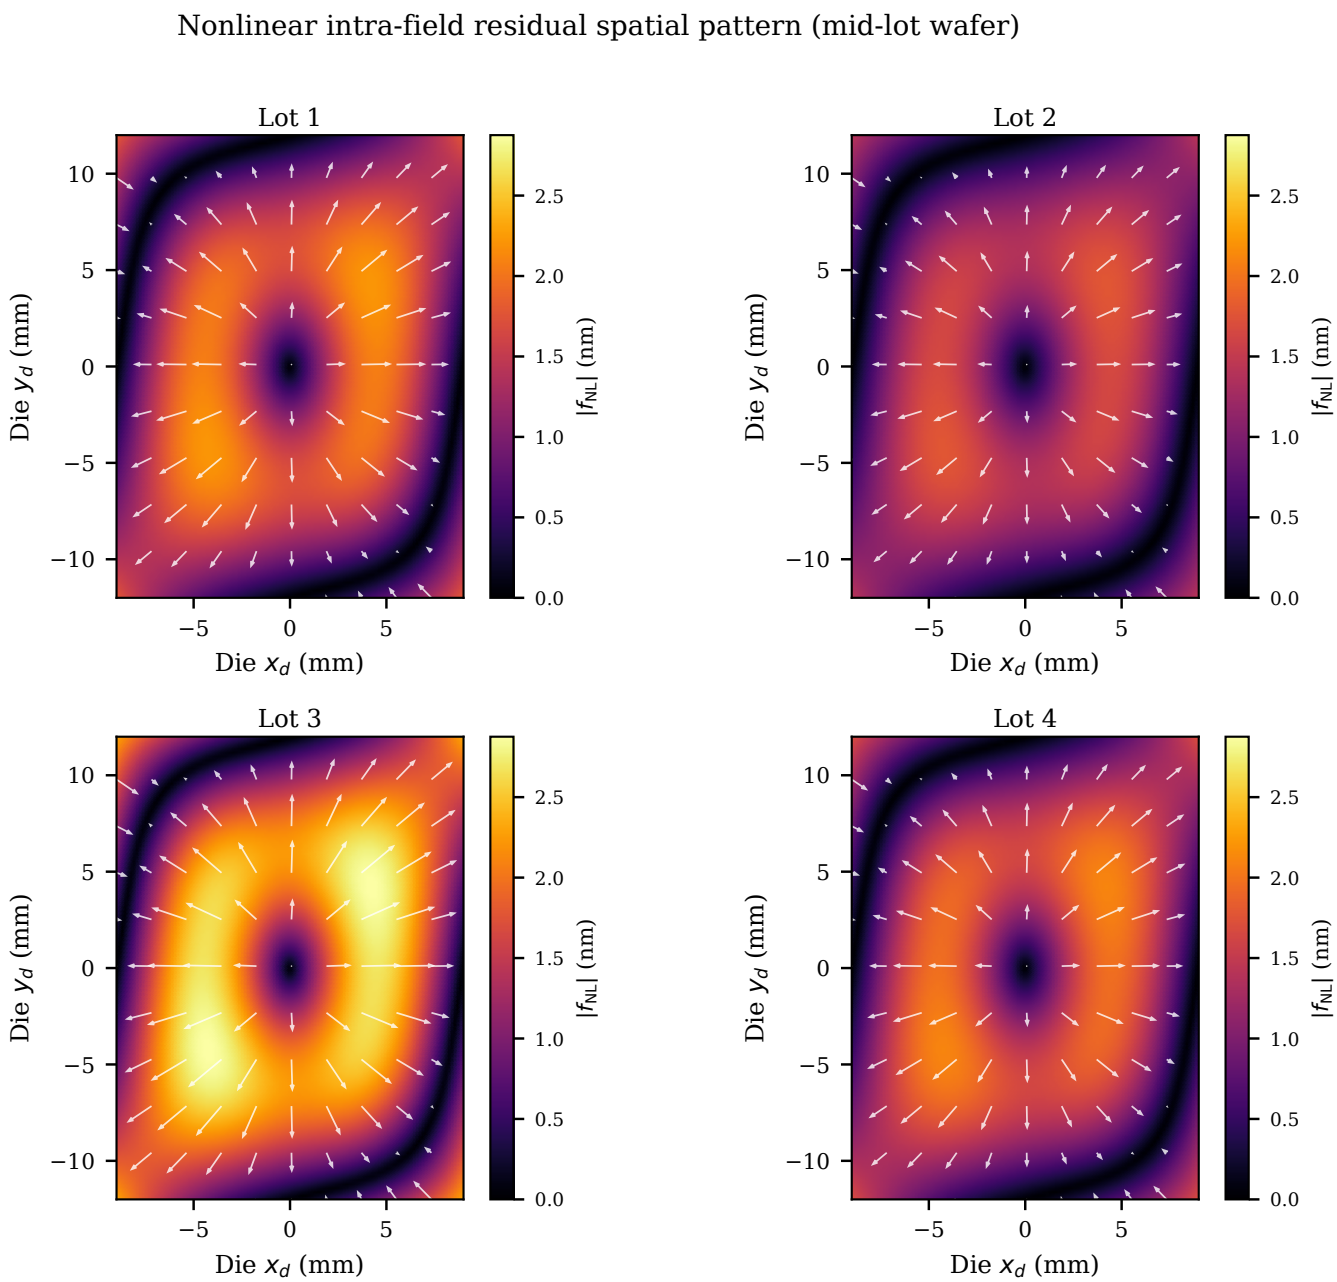

Fig. 4. Nonlinear intra-field residual f_{NL} evaluated on the die-coordinate grid for a mid-lot wafer ($w = 25$). Colour encodes magnitude $|f_{\mathrm{NL}}|$ (nm); arrows show direction. The pattern is identical for all 1 420 sites per wafer and grows by 30% from wafer 1 to wafer 50. Lot 3 has the largest amplitude ($A = 1.30$); Lot 2 the smallest ($A = 0.80$).

Data and Code Availability

The source code and dataset generator are publicly available on GitHub at https://github.com/arismagk/euv-ovl-syn. A permanently archived version of the dataset is deposited on Zenodo [19] at https://doi.org/10.5281/zenodo.19725620. The repository provides documentation and usage instructions to support full reproducibility.

Appendix

The complete dataset generator is the single file `src/generate_synthetic_data.py` in the companion repository. Executing

```
python src/generate_synthetic_data.py \
        --outdir dat_synthetic --seed 42
```

reproduces the 200 CSV files constituting the full dataset (`LotNWaferM.csv`, $N \in \{1, \ldots, 4\}$, $M \in \{1, \ldots, 50\}$) with bit-exact reproducibility across platforms (NumPy

default_rng, seed 42). Each file contains 1 420 rows with columns WaferCenterX_mm_, WaferCenterY_mm_, DieCenterX_mm_, DieCenterY_mm_, OverlayX, OverlayY. The schema is identical to the industrial APC metrology export format, enabling drop-in replacement of real data in existing correction pipelines.

REFERENCES

[1] H. J. Levinson, "Overlay control for high-NA EUV lithography," *J. Micro/Nanopatterning, Mater., Metrol.*, vol. 21, no. 2, p. 020901, 2022.

[2] J. van Schoot *et al.*, "High-NA EUV lithography scanner: performance overview," in *Proc. SPIE*, vol. 11177, 2019, p. 1117706.

[3] J. van Schoot, R. van Ballegoij, H. Butler, E. van Setten, G. Schiffelers, W. Bouman, S. van Gorp, K. Umstadter, J. Zimmermann, D. Golde, J. T. Neumann, and P. Graeupner, "Next step in Moore's law: high NA EUV system overview and first imaging and overlay performance," *J. Micro/Nanopatterning, Mater., Metrol.*, vol. 24, no. 1, p. 011009, 2024.

[4] M. Brink, "Use of vector fields to characterize reticle and wafer distortions," in *Proc. SPIE*, vol. 2197, 1994, pp. 167–177.

[5] M. A. van den Brink, H. Franken, S. Wittekoek, and T. van Goor, "Alignment and overlay performance of ASML step-and-scan systems," in *Proc. SPIE*, vol. 2726, 1996, pp. 734–751.

[6] J. Phillips *et al.*, "EUV overlay performance at advanced nodes," in *Proc. SPIE*, vol. 11328, 2020, p. 113280F.

[7] F. Houlihan *et al.*, "Lens heating effects in ArF immersion lithography," in *Proc. SPIE*, vol. 6520, 2007, p. 652011.

[8] A. Magklaras, P. Alefragis, C. Gogos, and A. Birbas, "Photolithography overlay errors simulation using Zernike based aberration modeling and intra-field variation," in *Proc. 10th South-East Europe Design Automation, Computer Engineering, Computer Networks and Social Media Conf. (SEEDA-CECNSM)*, Patras, Greece, 2025, pp. 1–6.

[9] M. Born and E. Wolf, *Principles of Optics*, 7th ed. Cambridge University Press, 1999.

[10] B. T. Smith, A. D. Chen, and M. J. Lee, "Overlay modeling and correction in advanced lithography," *IEEE Trans. Semicond. Manuf.*, vol. 31, no. 4, pp. 525–534, 2018.

[11] T. Abbe, J. Sturtevant, and T. Faure, "Machine learning applied to advanced lithography process control," in *Proc. SPIE*, vol. 10959, 2019, p. 109590D.

[12] S. Kim, J. Park, J. Choi *et al.*, "Deep learning-based overlay error prediction for advanced semiconductor manufacturing," in *Proc. SPIE*, vol. 11325, 2020, p. 113251A.

[13] M. Raissi, P. Perdikaris, and G. E. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *J. Comput. Phys.*, vol. 378, pp. 686–707, 2019.

[14] V. K. Pendyala, H. S. Sarjoughian, B. Potineni, and E. J. Yellig, "A benchmark time series dataset for semiconductor fabrication manufacturing constructed using component-based discrete-event simulation models," arXiv:2408.09307, 2024.

[15] A. Erdmann, *Optical and EUV Lithography: A Modeling Perspective*. SPIE Press, 2021.

[16] D. Rui, L. Zhang, Y. Wei, and Y. Su, "Method of high-order advanced lithography overlay correction to enhance the manufacturing performance of integrated circuits," *Nanoscale Adv.*, 2025.

[17] R. D. Guerrero, "Physics-constrained adaptive neural networks enable real-time semiconductor manufacturing optimization with minimal training data," arXiv:2511.12788, 2025.

[18] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016.

[19] A. Magklaras, G. Tsirogiannis, A. Birbas, C. Gogos, and P. Alefragis, "EUV-OVL-SYN: A scanner-informed synthetic overlay benchmark dataset for EUV lithography calibration," 2026. [Online]. Available: https://doi.org/10.5281/zenodo.19725620